

## IMPORTING RELEVANT

## LIBRARIES

We import necessary libraries including pandas for data handling, scikit-learn modules for machine learning tasks.

```
In [20]: import pandas as pd
import numpy as np
import seaborn as sn
from matplotlib import pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

## Loading Dataset and

## Inspecting

We load the Pima Indians Diabetes Dataset for inspecting into a pandas DataFrame

```
In [21]: df = pd.read_csv("C:/Users/Sanayak/Desktop/diabetes.csv")
df.head()
```

```
Out[21]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
In [22]: df.tail()
```

Out[22]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
<b>763</b>	10	101	76	48	180	32.9	
<b>764</b>	2	122	70	27	0	36.8	
<b>765</b>	5	121	72	23	112	26.2	
<b>766</b>	1	126	60	0	0	30.1	
<b>767</b>	1	93	70	31	0	30.4	

In [23]: `df.nunique()`

Out[23]:

Pregnancies	17
Glucose	136
BloodPressure	47
SkinThickness	51
Insulin	186
BMI	248
DiabetesPedigreeFunction	517
Age	52
Outcome	2

dtype: int64

In [24]: `df.describe()`

Out[24]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

## DATA CLEANSING AND DATA

### MANIPULATION

Our Dataset has zero(0) values for in columns

'Glucose','Bloodpressure','Skinthickness','Insulin' and 'Body Mass Index(BMI)'. We replace the zero(0) values with the median.

```
In [31]: df["BloodPressure"] = df["BloodPressure"].replace(0,72)
```

```
In [30]: df["Glucose"] = df["Glucose"].replace(0,117)
```

```
In [33]: df["SkinThickness"] = df["SkinThickness"].replace(0,23)
```

```
In [34]: df["Insulin"] = df["Insulin"].replace(0,30)
```

```
In [35]: df["BMI"] = df["BMI"].replace(0,32)
```

```
In [36]: df.describe()
```

```
Out[36]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Dia
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
<b>mean</b>	3.845052	121.656250	72.386719	27.334635	94.408854	32.450911	
<b>std</b>	3.369578	30.438286	12.096642	9.229014	105.695978	6.875366	
<b>min</b>	0.000000	44.000000	24.000000	7.000000	14.000000	18.200000	
<b>25%</b>	1.000000	99.750000	64.000000	23.000000	30.000000	27.500000	
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	31.000000	32.000000	
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	


## DATA VISUALISATION

```
In [38]: sn.pairplot(df)
```

```
Out[38]: <seaborn.axisgrid.PairGrid at 0x20ccba767e0>
```



```
In [41]: plt.hist(df.Age,bins=20)
plt.xlabel("Ages")
plt.ylabel("Frequency")
plt.title("Age Distribution")
plt.style.use("fivethirtyeight")
plt.show()
```



#### SPLITTING DATASET INTO FEATURE AND TARGET

VARIABLE

```
In [45]: X = df.drop('Outcome', axis=1)
         y = df['Outcome']
```

#### SPLITTING OUR DATASET INTO TRAINING SET

AND TESTING SET

```
In [47]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [ ]: #Standardizing features by removing the mean and scaling to unit variance
        scaler = StandardScaler()
        X_train = scaler.fit_transform(X_train)
        X_test = scaler.transform(X_test)
```

```
In [57]: # Calling LogisticRegression setting the maximum iteration to 400
        logreg = LogisticRegression(max_iter=400)
        logreg.fit(X_train, y_train)
        y_pred = logreg.predict(X_test)
        print(y_pred)
```

```
[0 0 0 0 0 0 0 0 1 1 0 1 0 0 0 0 0 0 1 1 0 0 1 0 1 1 0 0 0 0 1 1 1 1 1 1 1
0 0 1 0 1 1 0 0 1 1 0 0 1 0 1 1 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 0 1 1 0 0 0
0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 0 1 1 0 0 1 0 1 0 1 1 1 0 0 1 0 0 0
0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0
0 1 0 0 0 0]
```

```
In [58]: print(confusion_matrix(y_test, y_pred))
         print(classification_report(y_test, y_pred))
```

```
[[82 17]
 [20 35]]
```

	precision	recall	f1-score	support
0	0.80	0.83	0.82	99
1	0.67	0.64	0.65	55
accuracy			0.76	154
macro avg	0.74	0.73	0.74	154
weighted avg	0.76	0.76	0.76	154

TRAINED AND PREPARED BY VICTOR INIOBONG Economist | Mathematics Tutor| Data Scientist(In view)